

TrubiZone

Securing Critical Mobile Applications for Android Using ARM TrustZone

Tiago Luís de Oliveira Brito
Supervisor: Prof. Nuno Santos

Instituto Superior Técnico,
Av. Rovisco Pais, 1049-001 Lisboa - PORTUGAL
tiago.de.oliveira.brito@tecnico.ulisboa.pt

Abstract. With the ever-growing number of connected devices world-wide, and with a more conscious and sharing society, mobile devices are becoming interesting data banks. With such an impact in our lives, mobile developers must focus on developing secure and privacy aware applications to protect users and services. Today's mobile operating systems do not offer fully secure methods to support critical applications, such as e-Health, e-voting and e-money, by not taking advantage of secure hardware technology like TrustZone. This paper proposes a new model for the development of critical mobile applications and a system based on TrustZone implementing it.

1 Introduction

Mobile devices are becoming the predominant device. Actions previously performed by powerful desktop computers can now be easily replicated on mobile devices. A recent study [9] from early 2015, with key statistics for the U.S. market, even shows that mobile devices, such as smart-phones and tablets, dominate digital media time over the Personal Computer (PC), with the trend being to continue raising.

Present in our everyday personal and professional lives, mobile applications (apps) start to handle privacy and security-sensitive data. Most notably these apps are handling photos, health and banking information, location and general documentation. The growing role of mobile devices has the negative consequence of becoming an attractive target for attacks. Among the several mobile platforms, Android is the one which attracts much more malware attacks [21].

The health sector is an interesting market for attackers due to its information value. According to Reuters ¹, “medical identity theft is often not immediately identified by a patient or their provider, giving criminals years to milk such credentials. That makes medical data more valuable than credit cards”. This is way regulatory laws such as the Health Insurance Portability and Accountability

¹ <http://www.reuters.com/article/2014/09/24/us-cybersecurity-hospitals-idUSKCN0HJ21I20140924>

Act (HIPAA), which establishes standards for electronic health care transactions, are so important to follow when managing sensitive health data.

Valuable health care information is a motivation for criminals to perform critical attacks on hospital networks around the world, consequently leaking or stealing health records of millions of patients. In September 2014, a group of hackers attacked the network of University of California, Los Angeles (UCLA)'s Hospital accessing computers with sensitive records of 4.5 million people. According to Cable News Network (CNN) ², among the stolen records were the names, medical information, Social Security numbers, Medicare numbers, health plan IDs, birthdays and physical addresses of UCLA's patients.

In the mobile context, data is even more exposed and vulnerable due to the inherent portability of these devices, the sharing of information to third-party advertisers by device manufacturers or mobile app developers, unregulated management of sensitive medical information, specially because regulatory laws such as HIPAA do not account for the mobile sector, and because of security flaws on medical or consumer device software.

Parallel to the growth of the mobile app market, the number of e-Health mobile app, also known as mHealth (mobile health), available is increasing rapidly. In 2013, Research2Guidance [28] reported the existence of more than 97.000 mHealth apps across 62 app stores, with the top 10 apps generating up to 4 million free and 300.000 paid downloads per day. According to a report from MarketsAndMarkets [23], this market is expected to grow even further, from \$6.21 billion in revenue in 2013 to \$23.49 billion by 2018.

To protect sensitive data on mobile apps developers rely on mechanisms such as Digital Rights Management (DRM) [13], access control mechanisms, permission refinement, security Application Programming Interface (API), privacy enhancement systems and access control hooks, either from native Android ³, iOS ⁴ and Windows ⁵ or from extensions. These mechanisms rely on ad-hoc Operating System (OS) and application-level methodologies, which in most cases depend upon a very complex Trusted Computing Base (TCB) code, and do not fully enjoy the potential of the hardware of most modern smartphones, by not taking advantage of technology such as ARM's TrustZone.

Objetivos

Restrições

Especificação do sistema

Arquitetura

² <http://money.cnn.com/2015/07/17/technology/ucla-health-hack/>

³ <https://www.android.com/>

⁴ <http://www.apple.com/ios/>

⁵ <http://www.microsoft.com/en-us/windows>

Contribuições

The remainder of this document proceeds as follows. Section 2 highlights the related work on mobile health, mobile security and TrustZone technology. Section 3 highlights the architecture of TrubiZone.

2 Objectives (1pg)

Clearly explain the project objectives.

3 Related Work

This section presents the related work for this project and characterizes the main contributions of past works and how these contributions helped in the development of this project. This section is organized in three main parts: mHealth, mobile security and TrustZone.

The first part focuses on describing the attack surfaces of mHealth apps, the most common threats and their seriousness, present a few publicly available insecure mHealth apps as well as some compliance recommendations that app developers should follow to avoid unnecessary security risks when handling sensitive health information. The second part of this section describes the state-of-the-art of mobile security with a particular focus on the Android Operating System and how developers can build secure applications with a trusting OS and security mechanisms built upon the application layer. The third part of the related work describes TrustZone, a hardware technology available in most modern ARM Holdings (ARM) processors which supports execution of two isolated worlds, a hardware-protected secure world and a normal world. Besides describing the technology, this section also presents previous work developed using TrustZone and how these previous contributions may be helpful in achieving the main goals of this project.

3.1 mHealth

As discussed above, mobile devices are increasing in number at astonishing rates and with this growth the mobile market becomes cheap and accessible. This motivates the shift from mainframe systems, located in the facilities of healthcare providers, to apps on mobile devices as well as storage in shared cloud services. This accessibility also motivates the private sector in building more healthcare applications to support both patients and healthcare agents. Thus, the mobile health market is becoming a competitive market and one which is increasingly handling more sensitive data.

Kotz, David [20] defines a threat taxonomy for mHealth. In this paper the author categorizes threats into three main categories: *Identity Threats*, *Access Threats* and *Disclosure Threats*.

Identity threats are described as mis-use of patient identities and include cases where a patient may lose (or share) their identity credentials, allowing

malicious agents to access their Personal Health Record (PHR). *Insiders* (authorized PHR users, staff of the PHR organization, or staff of other mHealth support systems) may also use patient identities for medical fraud or for medical identity theft and *outsiders* may be able to observe patient identity or location from communications.

Access threats are described as unauthorized access to PHR and include cases where the patient, which controls the data, may allow broader-than-intended access or disclosure of information, *insiders* who may snoop or modify patient data, with intents other than improving the healthcare of the patient, and even *outsiders* which, by breaking into patient records, may leak or modify this data.

Disclosure threats include cases where an adversary captures network packets in order to obtain sensitive health data, this problem can be mitigated by using strong encryption methods, but even if the network traffic is encrypted it is possible to analyse the traffic to determine its characteristics [35]. The adversary may also use physical-layer or link-layer fingerprinting methods to identify the device type and, because the wireless medium is open, an active adversary may inject frames or may selectively interfere with (cause collisions with) wireless frames. These methods may enable the adversary to create a man-in-the-middle situation, to use link-layer fingerprinting methods, or to compromise the devices in a way that divulges their secrets.

This work by Kotz, David [20] helped in understanding what threats are inherent to mHealth systems and their development. An explicit set of rules for developing privacy-sensitive health applications was still needed to ease the development process. Avancha, Baxi and Kotz [1] filled the gap by surveying the literature, developing a privacy framework for mHealth and discussing the technologies that could support privacy-sensitive mHealth systems.

This survey considers essential accounting for privacy in the design and implementation of any mHealth system, given the sensitivity of the data collected. A developer may use the list of privacy properties provided by this article as a check-list that should be considered in any design. Furthermore, a set of questions is left open for researchers to improve the efficiency and effect of these properties. It is also mentioned that the privacy challenges identified by this article need to be addressed with urgency, because mHealth devices and systems are being deployed now, and retrofitting privacy protections is far more difficult than building them in from the start.

He, Dongjing, et al. [16] analyse several mHealth applications available in Android's app store contributing to the understanding of security and privacy risk on the Android platform.

In this paper, three studies were made with the following goals:

- **Study 1: What are the potential attack surfaces?**
- **Study 2: How widespread is the threat?**
- **Study 3: How serious is the threat?**

In the first study 160 apps are analysed to find evidence of threats. From analysing previous literature [36, 25, 2, 7, 8, 21], seven attack surfaces are deter-

Table 1. Description of attack surface (taken from He et al. [16])

Attack Surface	Description
Internet	Sensitive information is sent over the internet with unsecure protocols (e.g. HTTP), misconfigured HTTPS, etc.
Third Party	Sensitive information is stored in third party servers
Bluetooth	Sensitive information collected by Bluetooth-enabled health devices can be sniffed or injected
Logging	Sensitive information is put into system logs where it is not secured
SD Card Storage	Sensitive information is stored as unencrypted files on SD card, publicly accessible by any other app
Exported Components	Android app components, intended to be private, are set as exported, making them accessible by other apps
Side Channel	Sensitive information can be inferred by a malicious app with side channels, e.g. network package size, sequence, timing, etc.

mined to be in need of protection. These seven attack surfaces are shown in table 1.

In this study the authors also document that the 160 apps studied target two different audiences: 129 (81.65%) are for patients, 32 (20.25%) are for healthcare professionals and the remaining 3 (1.90%) are targeted for both. Most apps targeted for patients (60%) are in the Life Management category followed by apps that manage and synchronize user health information (PHR Management) which occupy nearly half (46.88%) of these apps. These numbers are a good indicator of what data is handled by most commercial mHealth apps available and motivate the choice made to build a PHR Management app as demo for TrubiZone.

A few examples of vulnerable applications are also revealed during this study. Regarding unencrypted information sent over the Internet, Doctor Online ⁶ (patients can talk to doctors online) and Recipes by Ingredients ⁷ send unencrypted sensitive information, including the user’s email and password, in clear text. With regard to logging sensitive information, the study pinpoints CVS/pharmacy ⁸, which logs the prescription refill details from user inputs, including name, email address, store number, and Rx number and also logs user login credentials in a debug log message. A malicious party can use this information to view user profiles, prescription history, which could support medical identity theft, and it is even possible to do online pharmacy shopping with user’s stored credit card information. Further applications are given as example for exposing sensitive information, Noom Weight Loss Coach ⁹ reveals user workout history by exposing its Content Providers to external apps, which means any app can access the

⁶ <https://play.google.com/store/apps/details?id=com.airpersons.airpersonsmobilehealth>

⁷ <https://play.google.com/store/apps/details?id=com.abMobile.recipebyingredient>

⁸ <https://play.google.com/store/apps/details?id=com.cvs.launchers.cvs>

⁹ <https://play.google.com/store/apps/details?id=com.wsl.noom>

exposed Content Providers without declaring any permission. Finally, this first study concludes with the example of sleep monitoring apps, such as SnoreClock¹⁰ and Sleep Talk Recorder¹¹, which store the sleep records of users as unencrypted audio files on external storage. With read permission for the SD card, as well as internet permission, a malicious app can read a user’s sleep recordings and even send them to remote servers.

In the second study, 27 apps of the top 1080 free apps from the Medical and Health & Fitness categories on Google Play were analysed according to their vulnerabilities. From this analysis, three attack surfaces are identified as the most important ones: *Internet*, *Third Party Services* and *Logging*. Only 7 of these 27 apps use the Internet to effectively send medical information over to remote servers. It is important to understand if the information sent over the Internet is protected. To achieve this, the authors captured network traffic and concluded that only 57.1% (4/7) of these apps use encrypted communication and the remaining 42.9% (3/7) use unencrypted communication to send sensitive health information. Among the unencrypted contents sent by these 3 apps were emails, usernames and passwords. This study also concludes that 85.7% (6/7) of these apps are hosted and store the recorded data on third party servers. This is an economical and scalable solution for mobile applications, but storing sensitive health records on third party servers can have serious implications, mostly due to app users not being aware that their data is being stored on third party servers and because these users are not able to tell if this data is stored encrypted in such a way that hosting companies do not have access to it.

Some health apps use Bluetooth devices to collect personal health information such as heart rate, respiration, pulse oximetry, electrocardiogram (ECG), blood pressure, body weight, body temperature, quality of sleep and exercise activities. Naveed et al. [25] show how a malicious app can stealthily collect user data from an Android device or spoof a device and inject fake data into the original device’s app, in what is called an external-device misbonding (DMB) attack. One of the 27 apps analysed in this study connects to external health sensors and uses default PIN code 0000, which makes it vulnerable to the DMB attack. Along with the logging, external storage, exported components and the other problems discussed above, which represent the explicit channels used for attacks, side channels can be exploited by a malicious party to infer sensitive information from apps, even when they are well-designed and implemented. This study mentions an example by Zhou et al. [36] where a correlation between network payload size and the disease condition a user selects on WebMD mobile¹². Network payload size is publicly accessible in Android, which represents a problem when such correlations can be made. Zhou et al. [36] mitigate the problem by modifying the Android kernel to enforce limitations on accessing Android public resources.

¹⁰ <https://play.google.com/store/apps/details?id=de.ralphsapps.snorecontrol>

¹¹ <https://play.google.com/store/apps/details?id=com.madinsweden.sleepstalk>

¹² <https://play.google.com/store/apps/details?id=com.webmd.android>




Fig. 1. Sensitive information distribution for study 3 (taken from He et al. [16])

In the third study, another 22 apps, which send information over the Internet, are randomly selected from the same top 1080 apps and used to understand the seriousness of the threat. The 22 apps are analysed to understand what information is effectively being sent over the Internet and the conclusion is that when used as intended these apps gather, store and transmit a variety of sensitive user data which includes at least personal profiles, health sensor data, lifestyle data, medical information browsing history and third-party app data (e.g. Facebook account information). Figure 1 shows the distribution of sensitive data in the 22 apps. Consequences of data breaches, information disclosure or tempering with sensitive health data depend on the type, sensitivity and volume of the data breached, but it is clear that profiling, medical identity theft and health-care decision-making errors are all possible. This is why the authors suggest the use of encryption for communication and storage, and encourage developers to create a set of security and privacy guidelines that offer a baseline for protection.

The work by He, Dongjing, et al. [16] helped in understanding what are the most common attack surfaces when considering mHealth applications on Android and it also helped assessing the security risks inherent to applications handling privacy sensitive information. This paper generally describes the state-of-the-art of security in commercial mobile health apps and pinpoints what should be fixed in order to build better and more secure application for the healthcare industry. It is clear that developers focus much more on building feature-full applications rather than secure apps and this is why it is important to build a framework which allows developers to completely focus on the features they want to provide without having to focus heavily on security. We believe that building a developer framework over ARM's TrustZone technology is a good solution

because with TrustZone is possible to fully enjoy the potential of the hardware already available in most mobile devices. This technology will be described in section 3.3.

3.2 Mobile Security Mechanisms

This section describes the security mechanisms available for the Android platform, either from native Android or as security extensions, and how these mechanisms can be used to solve some of the problems covered in the previous section. Work by Duarte, Nuno [14] and Costa, Miguel [11] extensively describe Android's security mechanisms and divide them in seven groups: *(i)* digital rights management, *(ii)* access control mechanisms, *(iii)* permission refinement, *(iv)* application communication monitoring, *(v)* privacy enhancement systems, *(vi)* access control hook APIs, and *(vii)* memory instrumentation. Although some of these systems may not be suited for the needs of the problem we aim to solve, this section aims to be a description of the state-of-the-art regarding security on the Android platform. Moreover, to the seven groups described, we suggest the addition of a new mechanism: *(viii)* embedded security systems. This last mechanism is not directly related to Android but, with a significant presence in most modern smartphones, its discussion is imperative for this project.

3.2.1 Digital Rights Management The first security mechanism analysed is Digital Rights Management. This specific access control technology allows data owners to restrict if and how their data is copied and also how that data is handled once transferred to another device. The DRM ecosystem is composed of the following entities:

- *User* - human user of the DRM Content
- *Content Issuer* - entity that delivers the content
- *Rights Issuer* - entity responsible for assigning permissions and constraints to DRM content
- *Rights Object* - XML document generated by a Rights Issuer expressing the restrictions associated to the content.
- *DRM Agent* - trusted entity responsible for enforcing permissions and constraints upon the DRM content

DRM can be used to protect proprietary software, hardware or general content. The most common application for DRM nowadays is music and video games, because these represent highly valuable intellectual property. To understand how DRM works in the context of mHealth one can suggest a simple example of a PHR mobile health application. In this scenario, the healthcare provider (e.g. a hospital) would be the *content issuer*, and it would use a *rights issuer* to assign the restrictions imposed upon the DRM content, which in this case would be the personal health record of a patient, when this content is transferred to the patient's device. When using DRM, the patient is limited to access the content through a *DRM Agent*.

The Open Mobile Alliance (OMA) developed a DRM standard [13] which defines the format of the content delivered to DRM Agents, as well as the way this content can be transferred from the Content Issuer to the DRM Agent. Android provides an extensible DRM framework, called Android DRM Framework¹³, allowing application developers to enable their apps to manage rights-protected content by complying with one of the supported DRM schemes (specific mechanisms, enforce by DRM Agents, to handle particular types of files).

This mechanism can be useful for protecting personal healthcare records by restricting the way data can be copied and handled. For this project however, DRM is not enough to ensure the confidentiality, integrity and availability of the sensitive data because it may fail to do so when the operating system is compromised with a malicious agent.

3.2.2 Access Control Mechanisms Access control mechanisms is a security model in which subjects (e.g. user, processes, threads, etc.) can perform actions on the system, namely on resources (e.g. files, sensors, etc.), typically called objects. Android follows a type of access control referenced as Discretionary Access Control (DAC), heritage from its Linux based kernel. In a DAC system the data owner is responsible for this data and thus determines who can access to it. In a Linux system, one can imagine a system administrator creating several files and allowing access to these files according to certain permissions. In this example, a subject with the specified permissions may access the file like its owner intended. In Android we can think of a similar example, since an application can create and store files in the filesystem, thus becoming the sole owner of these files, it can allow access to these files to any other application.

Although Android inherits the DAC from its Linux ancestry, most other resources in Android follow Mandatory Access Control (MAC) policies. In MAC, subjects are much more restricted in determining who has access to their resources. An example of this, from the physical security field, are the levels: confidential, secret and top secret. These labels are the only concept available to define the level of clearance of subjects or to classify data. When a subject attempts to access a classified piece of data a verification is done to assess if this subject's security level matches (or is above) that of the classified piece of data. Similarly, in Android once a subject attempts to access an object, it triggers a policy evaluation by the kernel, which assesses whether the access may be granted. The advantage of this strict system is its robustness, because subjects cannot override or modify the security policy. In Android, applications must specify in their manifests the permissions they require at runtime and after the installation neither applications nor users have any control over the access policies.

Because MAC is robust, several systems were created over the years to extend MAC's access control model to other Android resources. SEAndroid [32] solves problems related to resources complying with the DAC mechanism. The authors ported SELinux [27] to provide MAC at the kernel layer. The kernel

¹³ <http://developer.android.com/reference/android/drm/package-summary.html>

was then modified to support a new MAC policy (e.g., filesystem, IPC) and a new middleware layer (MMAC) was created to extend MAC to Android’s Binder IPC. TrustDroid [6] extends the MAC mechanism to all the platform’s resources (e.g., filesystem which complied with DAC) in order to isolate different domains’ sensitive information.

Even though access control mechanisms allow isolation between different subjects in a system, thus enforcing the confidentiality and integrity needed for this project, it still relies on the premise that the underlying middleware and kernel are secure. With this assumption, sensitive data may be at risk when the TCB is compromised. For this reason, access control mechanisms are not sufficient to ensure the security needs of this project.

3.2.3 Permission Refinement The Android permission mechanism is a very restrict system. At install time, a list of permissions an application specifies in its manifest file is presented to the user, which is forced into a binary decision, either granting all permissions forever or quitting the installation. This is an inflexible solution which makes it impossible for users to have full control of the permissions an application is effectively using at runtime. This inflexibility allows apps to use the device’s resources, authorized by the user at install time, whenever the app wishes without the user’s knowledge, possibly with malicious intent. Moreover, this permission mechanism allows for some problems such as component privilege escalation. An example of this is when an application, which profits from displaying in-app ads, requests access to the camera and to the Internet. In this case, the ad component has access to the camera, which it should not have because there is no need for it, and furthermore the component may even escalate its privileges by leaking video through the Internet.

Over the years many systems solved some of the problems inherent to permission refinement. APEX [24] modified Android’s permission model to allow users to specify the assignment of permissions to each app, both at install and runtime. Permission Tracker [19] allows users to be informed on how the permissions are used at runtime and offers the possibility of revoking these permissions. Furthermore, a user can specify which permissions are of interest so they can be notified of every permission access and decide whether to grant or deny that access. Compac [34] also restrains application access to certain permissions, but does so considering applications as groups of components and preventing component privilege escalation.

These systems improve upon the original Android permission model, but require manual configuration by the user. A more useful solution would be to use a context-aware system to handle the permissions at runtime. This way it may be possible to, without manual configuration, restrict permissions to all applications running along side a security sensitive application, thus isolating this application and avoiding possible leaks by shared resources. Several context-aware permission refinement systems exist. Trusted third parties can use CRePE [10] to enforce security policies on another devices whilst, for example, the employees’ mobile devices are inside the company, but have no privileges otherwise.

Similarly, MOSES [29] enforces domain isolation through the concept of security profiles. MOSES can switch security profiles based on pre-established conditions (e.g., GPS coordinates and time). When switching from a personal (less restrictive) profile to a professional (more restrictive) one, MOSES terminates and prevents the execution of applications not allowed to run in the new enforced profile. Additionally, MOSES leverages TaintDroid [15] (covered in Section 3.2.5) to prevent apps from one profile to access data belonging to another.

Both CRePE and MOSES suffer from a device control issue where a third party defines a policy that cannot be revoked by the user. Moreover, a user has no way to deny the enforcement of a third party policy. Duarte [14] and Costa [11] developed a system called TrUbi, built on top of ASM [17] (covered in Section 3.2.6), which solves the problems mentioned and presents a developer friendly framework for the development of sensitive permission refinement applications.

Even though TrUbi could be a good base for the development of this project, by allowing the development of applications which can control and enforce restrictions over less secure-sensitive apps, it suffers from the same limitations as the mechanisms discussed previously. When the underlying middleware and kernel are compromised, the security sensitive data may become accessible by a malicious agent.

3.2.4 Application Communication Monitoring Another problematic mechanism in Android is the way applications interact with each other. Some problems were already described in section 3.1 regarding vulnerable mobile health applications. Two attacks will be described in this section, followed by systems developed to mitigate such attacks on Android. The two attacks described are called *Confused Deputy Attacks* and *Collusion Attacks* and are represented in Figure 2. Confused deputy attacks basically consist of unprivileged applications taking advantage of other applications' publicly accessible APIs to perform malicious actions. Collusion attacks consist of an app, which might not have permission to perform an operation, still being able to perform it if there exists another app, belonging to the same developer, installed on the user's device with the permission to perform said operation. This happens because Android's permission system is based on UIDs.

Confused deputy attacks allow applications to use resources without explicitly specifying the necessary permission to do so. A simple example is of an app that exposes a service on its API allowing another app to receive photos directly from the device's camera. The second app, which did not specify the camera permission, can access this sensor without the user's knowledge. If this second app is malicious and has access to the Internet, then it is clear the severity of the problem which may lead to sensitive information being leaked. This problem generally happens because of poor or negligence application development.

A system called Saint [26] was created to specify which apps can access the public APIs of another app. This is done by allowing developers to assign appropriate security policies on their APIs. The identity of the caller apps is authenticated by the modified IPC mechanism. Other systems control this com-

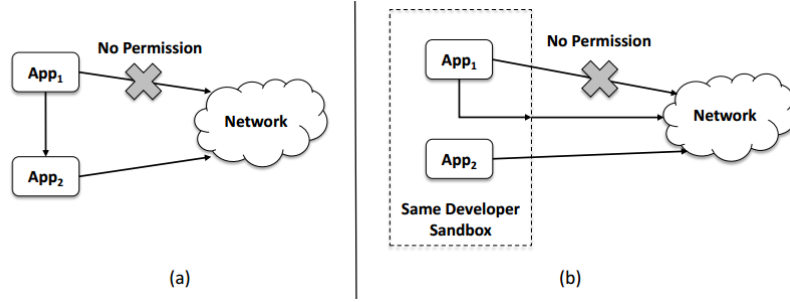


Fig. 2. Confused deputy attack (a); Collusion attack (b) - (taken from Duarte [14])

munication by extending Android’s Binder IPC mechanism, which is the main form of application inter-communication in Android. QUIRE [12] denies access to an API if in the message exchanged between apps, which contains the full call chain context, the source of the request does not have the necessary permission to access the corresponding data.

Collusion attacks are based on a malicious developer building a legitimate application with permissions such as microphone access and persuading the user to install another app with Internet permission. Both apps can collude to leak sensitive audio data through the Internet because Android’s permission system is based on UIDs, thus, although the first app does not explicitly request Internet access, this resource is accessible because the developer has Internet access through the first app. XManDroid [5] extends the Android permission model in order to support policies that could constrain the way apps interact with each other. This system prevents data leakage or other types of collusion attacks by developing a graph at runtime with a representation of the apps’ interactions, which is then used to regulate app inter-communication.

Although these systems assist the development of more secure mobile applications on the Android platform, by themselves these systems are not capable of building security sensitive apps and thus are not suitable for this project.

3.2.5 Privacy Enhancement Systems The most important concern for a privacy-sensitive mobile application is the way data is managed, specially with valuable data such as health records or banking information. Many systems have been developed to target different privacy problems in the Android environment.

Systems like MockDroid [4] and Zhou et al. [37] are extensions to Android’s data access control and prevent untrusted applications from accessing sensitive data by allowing users to manually specify application access rights over the system services providing access to such data. These systems even offer data shadowing¹⁴ to the unauthorized apps trying to access it. Other systems, like

¹⁴ Data shadow is the return of empty or incorrect information instead of the intended data.

IdentiDroid [31] focus on protecting the identity of the user by preventing applications from identifying the user and device. This is done by using an anonymous mobile device state to provide data shadowing and permission revocation techniques that disable the ability of apps to use systems services such as location or IMEI. This way if an app is trying to use the devices location a default location is returned instead of the real one or the application simply cannot access the location service.

Some systems use a different approach to solve the same problem. Dynamic taint analysis systems prevent data leakage by tainting data with a specific mark and then evaluate how this data is propagated through the system, If this data attempts to leave the system the user is alerted. TaintDroid [15] does this by modifying Android’s Dalvik VM environment and Binder IPC library to track the information within different applications’ processes. Although being an interesting system, TaintDroid suffers from limitations which include: *(i)* tracking a low number of data sources (mainly sensors), *(ii)* performance overheads not tolerable for most mobile environments, *(iii)* the existence of false positives leading to access control circumvention and *(iv)* the incapacity of analysing sensitive information leakage through covert channels.

Despite these limitations, TaintDroid was used by the research community as the building block for other privacy enhancement systems. Kynoid [30] extends TaintDroid to support a bigger number of information sources and to provide users a way to establish data access rights on data items and allows the to specify a list of acceptable IP destinations for each item. Another TaintDroid based system is AppFence [18] which provides a mechanism to block data made available by the user for on-device use only to be transmitted over the network, and allows for high level data to be substituted by shadowed data.

Although these systems aim at protecting users’ privacy they focus on detecting and preventing data leakage instead of providing a framework for developers to build privacy-sensitive applications. Even if a good framework for privacy-sensitive app development existed the limitation discussed before would persist. If the middleware or kernel layers became compromised then all systems dependent of these would be compromised as well.

3.2.6 Access Control Hook APIs A trend from previous work on mobile security is clearly noticeable: most security extensions require modifying and adding components to the kernel and middleware layers in order to implement new security models. Some frameworks have been built to ease the development process by allowing developers to easily create security models as ordinary apps whilst benefiting from a full callback system capable of notifying the security enforcement applications of access and requests to some of the apps’ resources of interest.

These frameworks comprise a set of hooks distributed among the kernel and middleware layers, which can be registered by a secure application. When a hook is activated it triggers a callback from the Hook API Module, which in turn is forwarded to the app for verification. The app then decides if operation

that triggered the activation of the hook may or may not proceed. The main advantage of frameworks such as Android Security Modules (ASM) [17] and Android Security Framework (ASF) [3] is the flexibility and freedom given to developers in choosing whatever resources to manage.

TrUbi [11, 14] is built upon ASM and allows for flexible system-wide resource restriction which may be used to isolate privacy sensitive apps by killing, freezing or revoke permissions to already running applications. The limitation of TrUbi is already discussed in section 3.2.3 and is based upon the fact that the TCB is comprised of all the Android kernel, Android middleware and extensions such as ASM.

3.2.7 Memory Instrumentation Memory instrumentation can be divided in two groups: static memory instrumentation and dynamic memory instrumentation. While static memory instrumentation changes already compiled byte-code, dynamic instrumentation patches running processes. In the remainder of this section the sole focus will be on dynamic memory instrumentation as it is an interesting mechanism to enforce a security model.

DeepDroid [33] relies on dynamic memory instrumentation to enforce fine-grained context-aware security policies for enterprise. This is done by patching several system services and tracing system calls.

These systems are better suited for the development of new security policies instead of being suited for the development of robust privacy-sensitive applications like the one we intend to build. Therefore, this security mechanism is not adequate for the needs of this project. Moreover, these systems still suffer from the same limitations discussed previously regarding compromised operating systems and large trusted computing base.

3.2.8 Embedded Security Systems Although developing new security mechanisms for mobile platforms like Android and iOS is important, mobile security is not limited to secure software. This section describes the embedded security architectures available on the market, along with its strengths and weaknesses. A whitepaper [22] by ARM states that security solutions for embedded applications fall into four categories, which will be explained in the remainder of this section.

External hardware security modules represent the classic security solution for embedded applications, which is the inclusion of a dedicated trusted element (e.g., a smartcard) that is outside of the main Socket-on-Chip (SoC). On one hand, the main advantage of this solution is that it allows for the encapsulation of the sensitive assets inside a physical device specially designed for robust security. On the other hand, the main disadvantage is that smartcards provide only secure processing and storage functions. This means that some operations (e.g., I/O) must rely on software running outside of the security perimeter to provide the desired features. An example where this happens is when a user interacts with the system by entering a Personal Identification Number (PIN).

This PIN must be managed by the less secure software outside of the smartcard, making it vulnerable to attacks.

Internal hardware security modules are, contrary to the previous solution, included within the SoC. These integrated modules are usually one of two forms: the first is a hardware block which manages cryptographic operations and key storage, and the second is a general purpose processing engine that is placed alongside the main processor, which uses custom hardware logic to prevent unauthorized access to sensitive resources.

This solution has the advantage of being cheaper, having the convenience of being an integrated part of the system and offering a performance improvement over external hardware security modules like smartcards. The disadvantage is that, like the previous solution, the resources protected by the security module will eventually need to be used outside the module itself, thus making these resources vulnerable. Another disadvantage is that this design requires a separate physical processor, typically less powerful than the main processor, to avoid sharing such an important resource with less critical modules. Extending or migrating systems built with this solution is also much harder. This is true because separation of resources must be implemented on proprietary hardware specific extensions within the SoC.

Software virtualization is another security mechanism which has been growing in popularity over the past years. In this software security mechanism a highly trusted management layer called hypervisor runs in a privileged mode of a general purpose processor. The hypervisor uses the Memory Management Unit (MMU) to separate several independent software platforms, running each one inside a virtual machine.

There are many advantages of using this solution. The first advantage is that a compact hypervisor can be thoroughly tested to ensure, with a high degree of certainty, that software running within one virtual machine cannot influence and attack the execution of others running in parallel. Another advantage is that there is no additional hardware requirement to implement a hypervisor, thus any processor with an MMU can be used to implement this security solution. Lastly, this solution is capable of supporting processing isolation between a secure environment and a full-blown rich operating system running in different virtual machines managed by the hypervisor. If a communication mechanism between virtual machines is supported by the hypervisor, then a secure pipeline is established between the two virtual machines.

The main problem with this approach is the isolation provided is restricted to the processor implementing the hypervisor. If all the other resources, mainly system buses, are not managed by the hypervisor, then the protection provided by the virtualization may be bypassed. Managing resources like Graphics Processing Units (GPUs) can be difficult to achieve without hindering the system's performance.

TrustZone is a hardware architecture that extends the security throughout the system design. The TrustZone architecture was designed to mitigate the problems discussed above. This is achieved by allowing any part of the system to be made secure instead of protecting only the assets inside a secure hardware block, thus enabling end-to-end security without exposing resources to a less critical platform.

To accomplish the goals of this project, a completely secure flow of execution is needed, this means that privacy-sensitive data must never leave the trusted security perimeter. A system which does not support secure I/O greatly limits the functionality offered by a PHR application, to the point where it becomes useless. For this reason we exclude both the external and internal hardware security modules as viable solutions for the problem.

On a different note, the virtualization solution seems promising, but with all the problems related to sharing resources such as GPUs and other system buses, this security module can be discarded as a solution. For an application to be secure, even when the rich OS is compromised, total isolation is needed, which is not easily achievable using virtualization.

TrustZone technology seems to be the most complete system of all, when considering embedded solutions, as it mitigates recurring problems from other mechanisms. Considering this, TrustZone is the most viable solution for our PHR applications, because it allows for a complete isolation between the privacy-sensitive data and all the other non-critical logic.

3.2.9 Summary After describing and assessing the available technology, one can argue which mechanisms better suit the development of critical mobile applications for Android, such as our PHR app. Some systems are clearly unsuitable for the goals of this project. Such mechanisms as application communication monitoring (section 3.2.4), access control hook APIs (section 3.2.6) or memory instrumentation (section 3.2.7) may be helpful in mitigating vulnerabilities caused by negligent development or for the development of new access control policies respectively, but are not sufficient to build privacy-sensitive applications. A summary of the main conclusions for the remaining systems studied is shown in table 2.

DRM by itself is capable of protecting data from disclosure and modification, hence it would be a great candidate for securing a privacy-sensitive application. Using access control mechanisms is possible to achieve data integrity and non-disclosure of sensitive data by restricting access to critical resources by untrusted applications, for example, by disabling Internet access to all untrusted apps one guarantees that sensitive data will not be leaked using this resource.

With permission refinement systems, in particular with TrUbi, it is possible to isolate the execution of a critical application from the remaining apps installed on the system, for example, the PHR app could be developed with the following premise. When the PHR app is running all the other apps are killed and all the resources blocked. With this the app could download the health records from

Table 2. Comparison between security mechanisms with and without a compromised Operating System (Android).

Security Mechanism	Trusted OS		Compromised OS	
	Confidentiality	Data Integrity	Confidentiality	Data Integrity
Digital Rights Management	✓	✓	✗	✗
Access Control Mechanisms	✓	✓	✗	✗
Permission Refinement	✓	✓	✗	✗
Privacy Enhancement Systems	✓	✓	✗	✗
Embedded Security Technology	✓	✓	✓	✓

the healthcare provider and show the data to the user. When the user exits the application, their data is encrypted with a key generated from a user password, and only then the resources are released for the other applications. The system was completely isolated during the whole process and the data is stored with encryption (and the key, because it is generated from a user password, is not stored on the system).

Although all of the above solutions are suitable for building a secure PHR application, they do not achieve the goals of this project, because all of these mechanisms suffer from the same limitation: both the middleware and the underlying kernel are considered as TCB. This means that if a malicious app is capable of compromising these layers, then the critical PHR application is compromised as well, and the data it protects may be accessible by the malicious party.

The only solution for the problem described is an hardware based security mechanism. As shown in table 2, the embedded security technology, particularly TrustZone, is the only solution which works either with a trusted OS or a compromised OS. For this reason, the next section will be focused on describing the state-of-the-art, architecture and functionality of the TrustZone technology, as it is the only security mechanism suitable for the needs of this project.

3.3 TrustZone

4 Architecture (2/3pgs)

Your proposed architecture. Can have lots of pictures and bullet points so it is easy to understand.

5 Evaluation (1/2pgs)

Explain how you are going to show your results (statistical data, cpu performance etc). Answer the following questions:

- Why is this solution going to be better than others.
- How am I going to defend that it is better.

6 Conclusions

Wrap up what you wrote.

References

1. Sasikanth Avancha, Amit Baxi, and David Kotz. Privacy in mobile technology for personal healthcare. *ACM Computing Surveys (CSUR)*, 45(1):3, 2012.
2. Adam J Aviv, Benjamin Sapp, Matt Blaze, and Jonathan M Smith. Practicality of accelerometer side channels on smartphones. In *Proceedings of the 28th Annual Computer Security Applications Conference*, pages 41–50. ACM, 2012.
3. Michael Backes, Sven Bugiel, Sebastian Gerling, and Philipp von Styp-Rekowsky. Android security framework: Enabling generic and extensible access control on android. *arXiv preprint arXiv:1404.1395*, 2014.
4. Alastair R Beresford, Andrew Rice, Nicholas Skehin, and Ripduman Sohan. Mockdroid: trading privacy for application functionality on smartphones. In *Proceedings of the 12th Workshop on Mobile Computing Systems and Applications*, pages 49–54. ACM, 2011.
5. Sven Bugiel, Lucas Davi, Alexandra Dmitrienko, Thomas Fischer, and Ahmad-Reza Sadeghi. Xmandroid: A new android evolution to mitigate privilege escalation attacks. *Technische Universität Darmstadt, Technical Report TR-2011-04*, 2011.
6. Sven Bugiel, Lucas Davi, Alexandra Dmitrienko, Stephan Heuser, Ahmad-Reza Sadeghi, and Bhargava Shastri. Practical and lightweight domain isolation on android. In *Proceedings of the 1st ACM workshop on Security and privacy in smartphones and mobile devices*, pages 51–62. ACM, 2011.
7. Liang Cai and Hao Chen. *On the practicality of motion based keystroke inference attack*. Springer, 2012.
8. Erika Chin, Adrienne Porter Felt, Kate Greenwood, and David Wagner. Analyzing inter-application communication in android. In *Proceedings of the 9th international conference on Mobile systems, applications, and services*, pages 239–252. ACM, 2011.
9. Comscore. Digital Future in Focus. Technical report, Comscore Inc., 03 2015. <http://www.comscore.com/Insights/Presentations-and-Whitepapers/2015/2015-US-Digital-Future-in-Focus>.
10. Mauro Conti, Vu Thien Nga Nguyen, and Bruno Crispo. Crepe: Context-related policy enforcement for android. In *Information Security*, pages 331–345. Springer, 2011.
11. Miguel Costa. SecMos - Mobile Operating System Security. Master’s thesis, Instituto Superior Técnico, Portugal, 2015.
12. Michael Dietz, Shashi Shekhar, Yuliy Pisetsky, Anhei Shu, and Dan S Wallach. Quire: Lightweight provenance for smart phone operating systems. In *USENIX Security Symposium*, page 24, 2011.
13. OMA DRM. Open mobile alliance digital rights management.(2010). Retrieved May, 2, 2011.
14. Nuno Duarte. On The Effectiveness of Trust Leases in Securing Mobile Applications. Master’s thesis, Instituto Superior Técnico, Portugal, 2015.
15. William Enck, Peter Gilbert, Seungyeop Han, Vasant Tendulkar, Byung-Gon Chun, Landon P Cox, Jaeyeon Jung, Patrick McDaniel, and Anmol N Sheth. Taintdroid: an information-flow tracking system for realtime privacy monitoring on smartphones. *ACM Transactions on Computer Systems (TOCS)*, 32(2):5, 2014.

16. Dongjing He, Muhammad Naveed, Carl A Gunter, and Klara Nahrstedt. Security concerns in android mhealth apps. In *AMIA Annual Symposium Proceedings*, volume 2014, page 645. American Medical Informatics Association, 2014.
17. Stephan Heuser, Adwait Nadkarni, William Enck, and Ahmad-Reza Sadeghi. Asm: A programmable interface for extending android security. In *Proc. 23rd USENIX Security Symposium (SEC'14)*, 2014.
18. Peter Hornyack, Seungyeop Han, Jaeyeon Jung, Stuart Schechter, and David Wetherall. These aren't the droids you're looking for: retrofitting android to protect data from imperious applications. In *Proceedings of the 18th ACM conference on Computer and communications security*, pages 639–652. ACM, 2011.
19. Michael Kern and Johannes Sametinger. Permission tracking in android. In *The Sixth International Conference on Mobile Ubiquitous Computing, Systems, Services and Technologies UBIComm*, pages 148–155, 2012.
20. David Kotz. A threat taxonomy for mhealth privacy. In *COMSNETS*, pages 1–6, 2011.
21. F-Secure Labs. Mobile Threat Report. Technical report, F-Secure Labs., 03 2014. https://www.f-secure.com/documents/996508/1030743/Mobile_Threat_Report_Q1_2014.pdf.
22. ARM Limited. ARM Security Technology - Building a Secure System using TrustZone Technology. Technical report, ARM Limited, 04 2009. <http://infocenter.arm.com/help/topic/com.arm.doc.prd29-genc-009492c/PRD29-GENC-009492C-trustzone-security-whitepaper.pdf>.
23. MarketsAndMarkets. Mobile health apps & solutions market by connected devices (cardiac monitoring, diabetes management devices), health apps (exercise, weight loss, women's health, sleep and meditation), medical apps (medical reference) – global trends & forecast to 2018. Technical report, MarketsAndMarkets, 09 2013. <http://marketsandmarkets.com/>.
24. Mohammad Nauman, Sohail Khan, and Xinwen Zhang. Apex: extending android permission model and enforcement with user-defined runtime constraints. In *Proceedings of the 5th ACM Symposium on Information, Computer and Communications Security*, pages 328–332. ACM, 2010.
25. Muhammad Naveed, Xiaoyong Zhou, Soteris Demetriou, XiaoFeng Wang, and Carl A Gunter. Inside job: Understanding and mitigating the threat of external device mis-bonding on android. In *Proceedings of the 21st Annual Network and Distributed System Security Symposium (NDSS)*, pages 23–26, 2014.
26. Machigar Ongtang, Stephen McLaughlin, William Enck, and Patrick McDaniel. Semantically rich application-centric security in android. *Security and Communication Networks*, 5(6):658–673, 2012.
27. NSA Peter Loscocco. Integrating flexible support for security policies into the linux operating system. In *Proceedings of the FREENIX Track:... USENIX Annual Technical Conference*, page 29. The Association, 2001.
28. Research2Guidance. Mobile Health Market Report 2013-2017. Technical report, Research2Guidance, 03 2013. <http://research2guidance.com/>.
29. Giovanni Russello, Mauro Conti, Bruno Crispo, and Earlene Fernandes. Moses: supporting operation modes on smartphones. In *Proceedings of the 17th ACM symposium on Access Control Models and Technologies*, pages 3–12. ACM, 2012.
30. Daniel Schreckling, Johannes Köstler, and Matthias Schaff. Kynoid: real-time enforcement of fine-grained, user-defined, and data-centric security policies for android. *Information Security Technical Report*, 17(3):71–80, 2013.
31. Bilal Shebaro, Oluwatosin Ogunwuyi, Daniele Midi, and Elisa Bertino. Identidroid: Android can finally wear its anonymous suit. 2014.

32. Stephen Smalley and Robert Craig. Security enhanced (se) android: Bringing flexible mac to android. In *NDSS*, volume 310, pages 20–38, 2013.
33. Xueqiang Wang, Kun Sun, Yuewu Wang, and Jiwu Jing. Deepdroid: Dynamically enforcing enterprise policy on android devices. In *Proc. 22nd Annual Network and Distributed System Security Symposium (NDSS’15)*. The Internet Society, 2015.
34. Yifei Wang, Srinivas Hariharan, Chenxi Zhao, Jiaming Liu, and Wenliang Du. Compac: Enforce component-level access control in android. In *Proceedings of the 4th ACM Conference on Data and Application Security and Privacy*, pages 25–36. ACM, 2014.
35. Charles V Wright, Fabian Monrose, and Gerald M Masson. On inferring application protocol behaviors in encrypted network traffic. *The Journal of Machine Learning Research*, 7:2745–2769, 2006.
36. Xiaoyong Zhou, Soteris Demetriou, Dongjing He, Muhammad Naveed, Xiaorui Pan, XiaoFeng Wang, Carl A Gunter, and Klara Nahrstedt. Identity, location, disease and more: Inferring your secrets from android public resources. In *Proceedings of the 2013 ACM SIGSAC conference on Computer & communications security*, pages 1017–1028. ACM, 2013.
37. Yajin Zhou, Xinwen Zhang, Xuxian Jiang, and Vincent W Freeh. Taming information-stealing smartphone applications (on android). In *Trust and Trustworthy Computing*, pages 93–107. Springer, 2011.